

# Car Price Prediction

```
In [ ]: #Importing the Libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [ ]: #Loading the dataset
df = pd.read_csv('car data.csv')
df.head()
```

```
Out[ ]:
```

	Car_Name	Year	Selling_Price	Present_Price	Driven_kms	Fuel_Type	Selling_type	T
0	ritz	2014	3.35	5.59	27000	Petrol	Dealer	
1	sx4	2013	4.75	9.54	43000	Diesel	Dealer	
2	ciaz	2017	7.25	9.85	6900	Petrol	Dealer	
3	wagon r	2011	2.85	4.15	5200	Petrol	Dealer	
4	swift	2014	4.60	6.87	42450	Diesel	Dealer	

## Data Preprocessing

```
In [ ]: #Checking the shape of the dataset
df.shape
```

```
Out[ ]: (301, 9)
```

```
In [ ]: #Checking for the null values
df.isnull().sum()
```

```
Out[ ]: Car_Name      0
Year            0
Selling_Price    0
Present_Price    0
Driven_kms       0
Fuel_Type        0
Selling_type     0
Transmission     0
Owner            0
dtype: int64
```

```
In [ ]: #checking for duplicate values
df.duplicated().sum()
```

```
Out[ ]: 2
```

```
In [ ]: #removing the duplicate values
df.drop_duplicates(inplace=True)
```

```
In [ ]: #Checking the datatypes of the columns
df.dtypes
```

```
Out[ ]: Car_Name      object
Year          int64
Selling_Price  float64
Present_Price  float64
Driven_kms     int64
Fuel_Type      object
Selling_type   object
Transmission   object
Owner          int64
dtype: object
```

```
In [ ]: #Checking the unique values in the categorical columns
df.nunique()
```

```
Out[ ]: Car_Name      98
Year          16
Selling_Price  156
Present_Price  148
Driven_kms     206
Fuel_Type      3
Selling_type   2
Transmission   2
Owner          3
dtype: int64
```

Descriptive Statistics

```
In [ ]: df.describe()
```

```
Out[ ]:
```

	Year	Selling_Price	Present_Price	Driven_kms	Owner
<b>count</b>	299.000000	299.000000	299.000000	299.000000	299.000000
<b>mean</b>	2013.615385	4.589632	7.541037	36916.752508	0.043478
<b>std</b>	2.896868	4.984240	8.566332	39015.170352	0.248720
<b>min</b>	2003.000000	0.100000	0.320000	500.000000	0.000000
<b>25%</b>	2012.000000	0.850000	1.200000	15000.000000	0.000000
<b>50%</b>	2014.000000	3.510000	6.100000	32000.000000	0.000000
<b>75%</b>	2016.000000	6.000000	9.840000	48883.500000	0.000000
<b>max</b>	2018.000000	35.000000	92.600000	500000.000000	3.000000

## Exploratory Data Analysis

### Top 10 Cars by Price

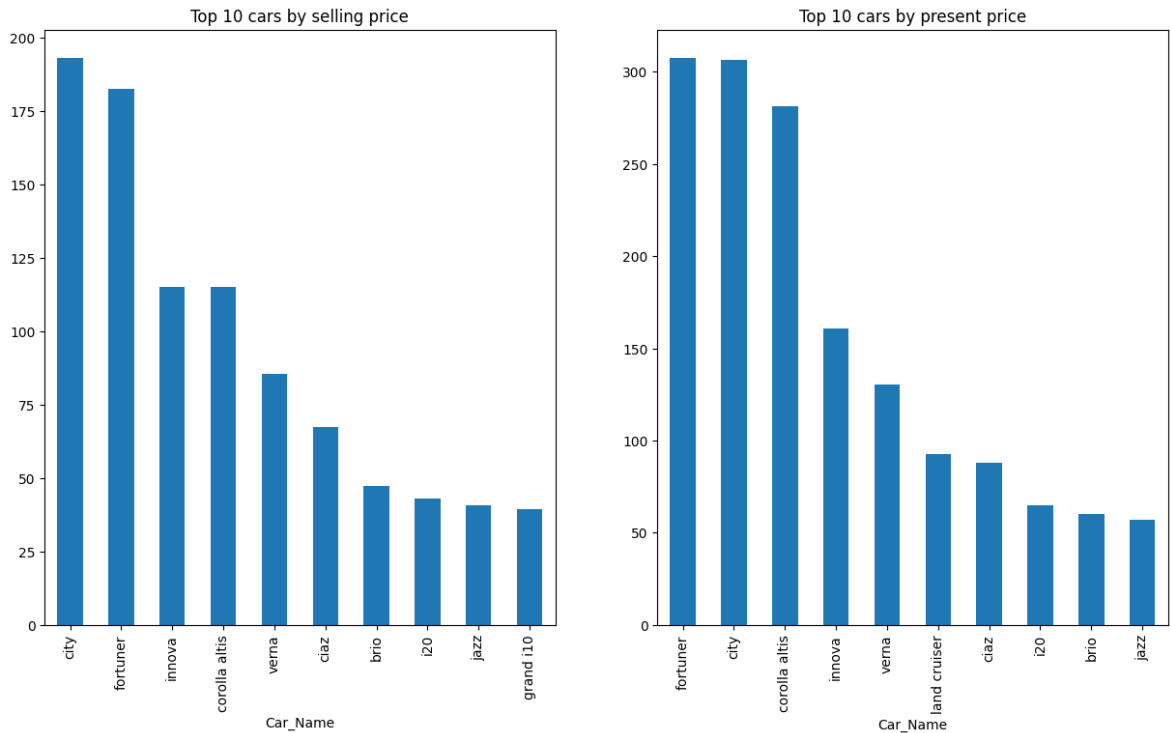
```
In [ ]: fig, ax = plt.subplots(1,2,figsize=(15,8))

#top 10 cars by selling price
```

```
df.groupby('Car_Name')['Selling_Price'].sum().sort_values(ascending=False).head(
ax[0].set_title('Top 10 cars by selling price')

#top 10 cars by present price
df.groupby('Car_Name')['Present_Price'].sum().sort_values(ascending=False).head(
ax[1].set_title('Top 10 cars by present price')
```

Out[ ]: Text(0.5, 1.0, 'Top 10 cars by present price')



## Year vs Price

```
In [ ]: fig, ax = plt.subplots(1,2,figsize=(15,8))

#variation of selling price with year
sns.lineplot(x='Year', y='Selling_Price', data=df, ax=ax[0]).set_title('Variatio

#variation of present price with year
sns.lineplot(x='Year', y='Present_Price', data=df, ax=ax[1]).set_title('Variatio
```

Out[ ]: Text(0.5, 1.0, 'Variation of present price with year')



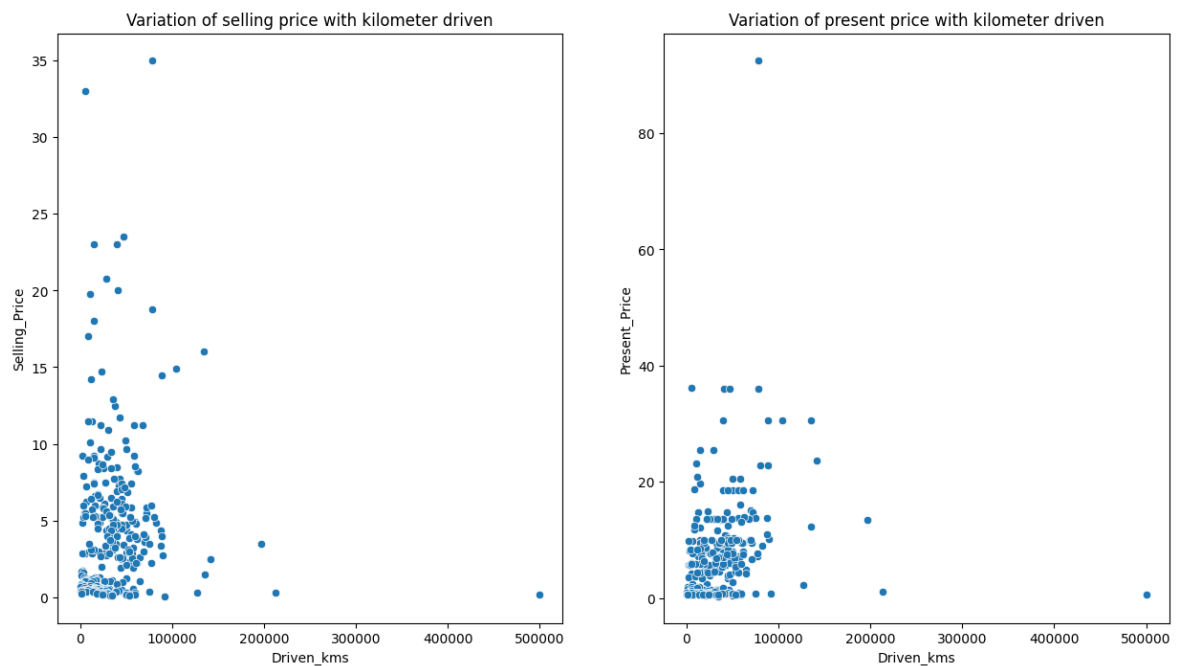
## Kilometers driven vs Price

```
In [ ]: fig, ax = plt.subplots(1,2,figsize=(15,8))

#variation of selling price with kilometer driven
sns.scatterplot(x='Driven_kms', y='Selling_Price', data=df, ax=ax[0]).set_title('Variation of selling price with kilometer driven')

#variation of present price with kilometer driven
sns.scatterplot(x='Driven_kms', y='Present_Price', data=df, ax=ax[1]).set_title('Variation of present price with kilometer driven')
```

```
Out[ ]: Text(0.5, 1.0, 'Variation of present price with kilometer driven')
```



## Fuel Type vs Price

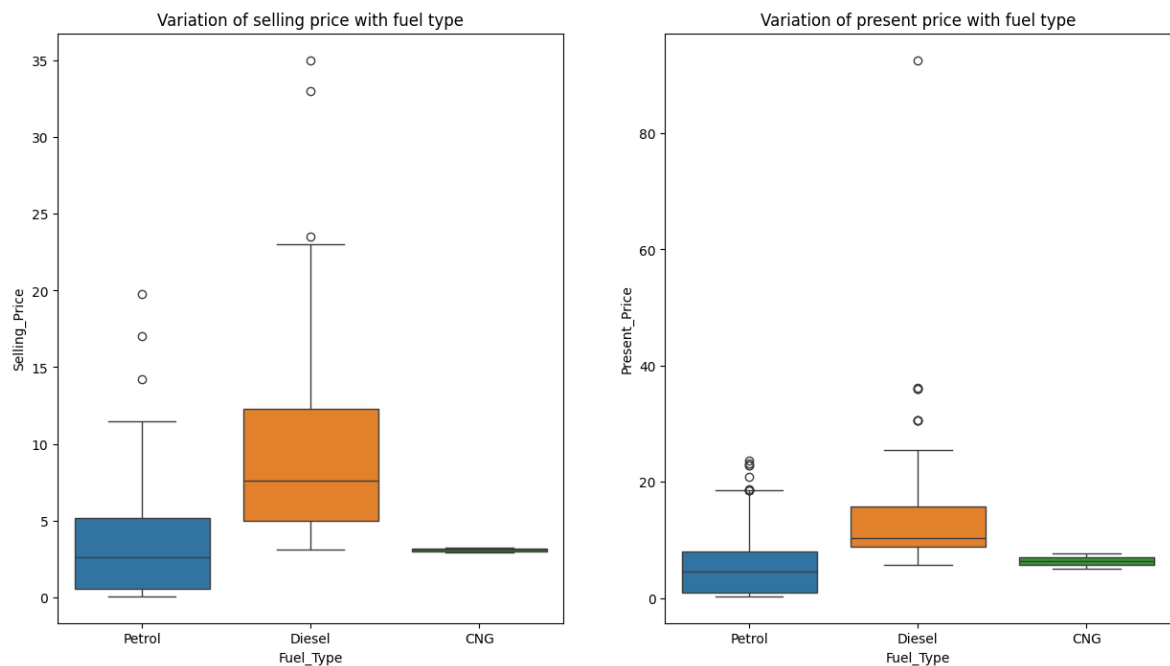
```
In [ ]: fig, ax = plt.subplots(1,2,figsize=(15,8))

#variation of selling price with fuel type
```

```
sns.boxplot(x='Fuel_Type', y='Selling_Price', data=df, ax=ax[0], hue = 'Fuel_Type')

#variation of present price with fuel type
sns.boxplot(x='Fuel_Type', y='Present_Price', data=df, ax=ax[1], hue = 'Fuel_Type')
```

Out[ ]: Text(0.5, 1.0, 'Variation of present price with fuel type')



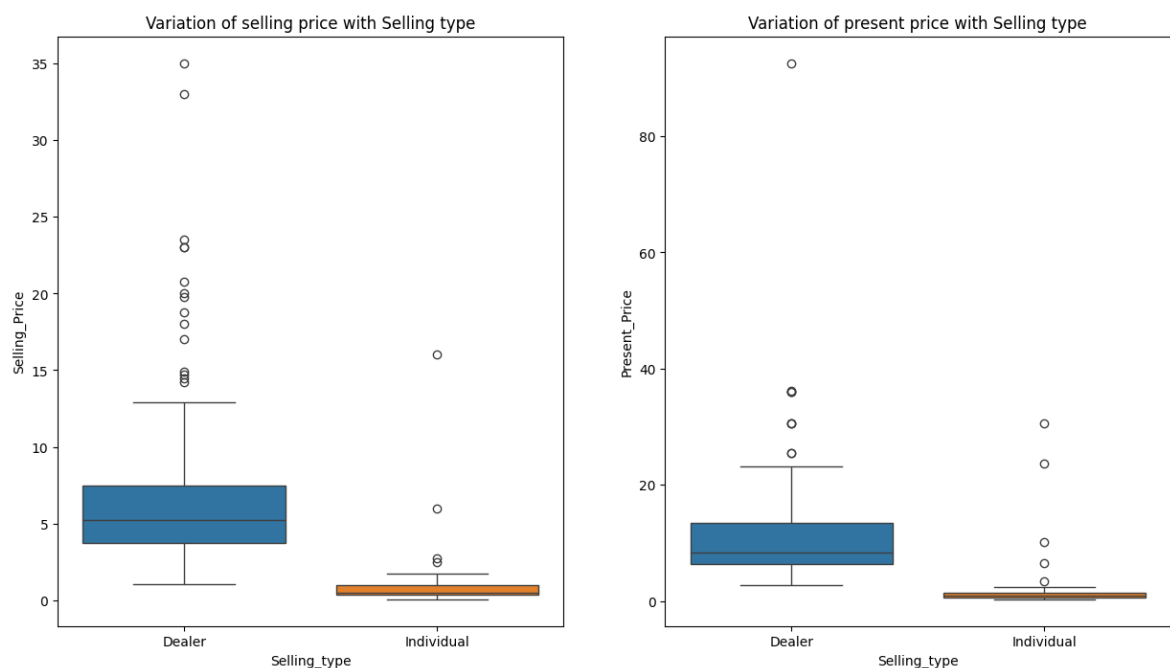
## Seller Type vs Price

```
In [ ]: fig, ax = plt.subplots(1,2,figsize=(15,8))

#variation of selling price with Selling_type
sns.boxplot(x='Selling_type', y='Selling_Price', data=df, ax=ax[0], hue = 'Selling_type')

#variation of present price with Selling_type
sns.boxplot(x='Selling_type', y='Present_Price', data=df, ax=ax[1], hue = 'Selling_type')
```

Out[ ]: Text(0.5, 1.0, 'Variation of present price with Selling type')



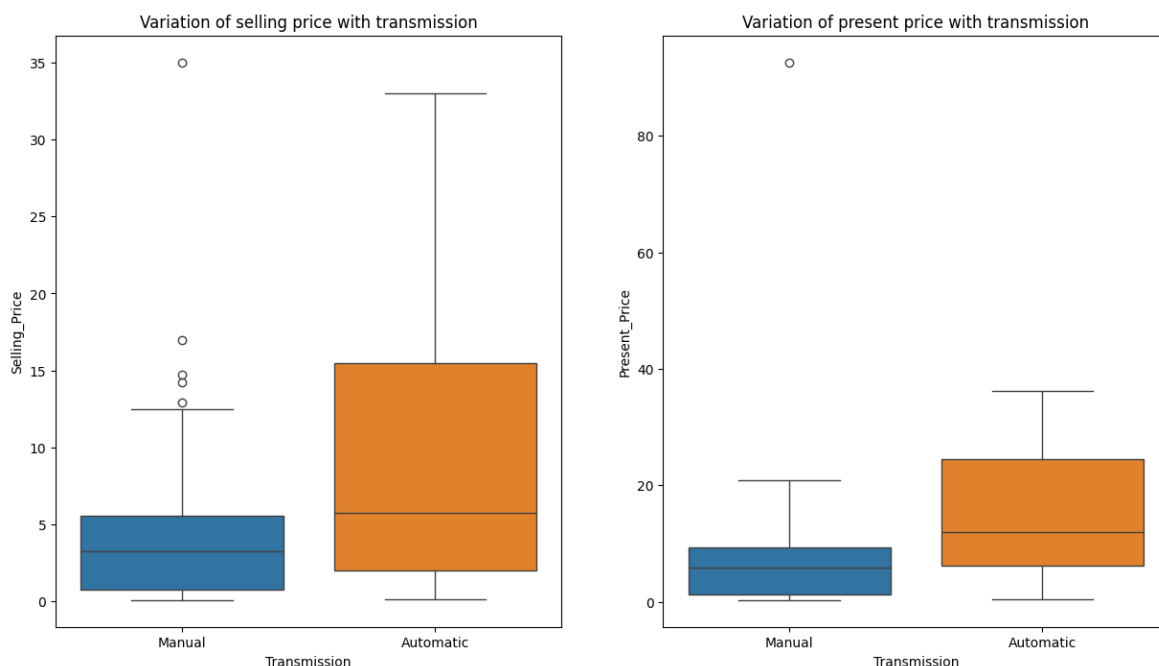
## Transmission vs Price

```
In [ ]: fig, ax = plt.subplots(1,2,figsize=(15,8))

#variation of selling price with transmission
sns.boxplot(x='Transmission', y='Selling_Price', data=df, ax=ax[0], hue = 'Trans

#variation of present price with transmission
sns.boxplot(x='Transmission', y='Present_Price', data=df, ax=ax[1], hue = 'Trans
```

```
Out[ ]: Text(0.5, 1.0, 'Variation of present price with transmission')
```




## Owner vs Price

```
In [ ]: fig, ax = plt.subplots(1,2,figsize=(15,8))

#variation of selling price with owner
sns.boxplot(x='Owner', y='Selling_Price', data=df, ax=ax[0], hue = 'Owner').set_

#variation of present price with owner
sns.boxplot(x='Owner', y='Present_Price', data=df, ax=ax[1], hue = 'Owner').set_
```

```
Out[ ]: Text(0.5, 1.0, 'Variation of present price with owner')
```



## Data Preprocessing 2

The Car\_Name has so many unique values, which make it difficult to train the machine learning model. So, we will drop this column, in order to reduce the parameters of the model.

```
In [ ]: #dropping the car name column
df.drop('Car_Name', axis=1, inplace=True)
```

## Label Encoding

```
In [ ]: from sklearn.preprocessing import LabelEncoder

#columns for label encoding
cols = df.select_dtypes(include='object').columns

#label encoding object
le = LabelEncoder()


#applying label encoding
for i in cols:
    df[i] = le.fit_transform(df[i])
    print(i, df[i].unique())
```

```
Fuel_Type [2 1 0]
Selling_type [0 1]
Transmission [1 0]
```

## Coorelation Matrix Heatmap

```
In [ ]: sns.heatmap(df.corr(), annot=True)
```

```
Out[ ]: <Axes: >
```



## Train Test Split

```
In [ ]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(df.drop('Selling_Price', axis=1), df['Selling_Price'], test_size=0.2, random_state=42)
```

## Model Building

I will using the following algorithms to build the models:

- Desicion Tree Regressor
- Random Forest Regressor

## Decision Tree Regressor

```
In [ ]: from sklearn.tree import DecisionTreeRegressor

#Decision Tree Regressor object
dt = DecisionTreeRegressor()
```

## Hyperparameter Tuning using GridSearchCV

```
In [ ]: from sklearn.model_selection import GridSearchCV

#parameters for grid search
para = {
```



```

    'max_depth': [2,4,6,8,10],
    'min_samples_split': [2,4,6,8,10],
    'min_samples_leaf': [2,4,6,8,10],
    'random_state': [0,42]
}

#Grid Search object
grid = GridSearchCV(dt, para, cv=5, n_jobs=-1, verbose=2, scoring='accuracy')

#fitting the grid search object
grid.fit(X_train, y_train)

#best parameters
print(grid.best_params_)

```

Fitting 5 folds for each of 250 candidates, totalling 1250 fits

```
{'max_depth': 2, 'min_samples_leaf': 2, 'min_samples_split': 2, 'random_state': 0}
```

```

In [ ]: #decision tree regressor object with best parameters
dt = DecisionTreeRegressor(max_depth=2, min_samples_leaf=2, min_samples_split=2,

#fitting the decision tree regressor object
dt.fit(X_train, y_train)

#training accuracy
print(dt.score(X_train, y_train))

#prediction on test data
dt_pred = dt.predict(X_test)

```

0.8127778422312646

## Random Forest Regressor

```

In [ ]: from sklearn.ensemble import RandomForestRegressor

#Random Forest Regressor object
rf = RandomForestRegressor()

```

## Hyperparameter Tuning using GridSearchCV

```

In [ ]: from sklearn.model_selection import GridSearchCV

#parameters for grid search
para = {
    'n_estimators': [100,200],
    'max_depth': [2,4,6],
    'min_samples_split': [2,4,6],
    'min_samples_leaf': [2,4,6],
    'random_state': [0,42]
}

#Grid Search object
grid = GridSearchCV(rf, para, cv=5, n_jobs=-1, verbose=2, scoring='accuracy')

#fitting the grid search object
grid.fit(X_train, y_train)

```

```
#best parameters  
print(grid.best_params_)
```

Fitting 5 folds for each of 108 candidates, totalling 540 fits  
{'max\_depth': 2, 'min\_samples\_leaf': 2, 'min\_samples\_split': 2, 'n\_estimators': 100, 'random\_state': 0}

```
In [ ]: #random forest regressor object with best parameters  
rf = RandomForestRegressor(max_depth=6, min_samples_leaf=2, min_samples_split=2,  
  
#fitting the random forest regressor object  
rf.fit(X_train, y_train)  
  
#training accuracy  
print(rf.score(X_train, y_train))  
  
#prediction on test data  
y_pred = rf.predict(X_test)
```


0.9505984366262189

## Model Evaluation

### Distribution Plot

```
In [ ]: ax = sns.distplot(y_test, color='r', label='Actual', hist = False)  
#decision tree regressor  
sns.distplot(dt_pred, color='b', label='Predicted', ax=ax, hist = False)  
#random forest regressor  
sns.distplot(y_pred, color='g', label='Predicted', ax=ax, hist = False)
```

```
Out[ ]: <Axes: xlabel='Selling_Price', ylabel='Density'>
```



## Model Metrics

```
In [ ]: from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error
```

```
#decision tree regressor
print('Decision Tree Regressor')
print('Mean Squared Error: ', mean_squared_error(y_test, dt_pred))
print('Mean Absolute Error: ', mean_absolute_error(y_test, dt_pred))
print('R2 Score: ', r2_score(y_test, dt_pred))
print('\n')
#random forest regressor
print('Random Forest Regressor')
print('Mean Squared Error: ', mean_squared_error(y_test, y_pred))
print('Mean Absolute Error: ', mean_absolute_error(y_test, y_pred))
print('R2 Score: ', r2_score(y_test, y_pred))
```

Decision Tree Regressor  
Mean Squared Error: 15.857485535955778  
Mean Absolute Error: 2.2808490157116004  
R2 Score: 0.3847311450436749

Random Forest Regressor  
Mean Squared Error: 10.159339309688956  
Mean Absolute Error: 1.395447827308639  
R2 Score: 0.6058186494944606